

Alex Chen

San Francisco, CA | alex.chen@techmail.dev | +1 (415) 555-0192 | linkedin.com/in/alexchen-ai

PROFESSIONAL SUMMARY

Staff / Senior Full-Stack AI Engineer with 6+ years of experience designing and deploying enterprise GenAI systems, high-throughput FastAPI microservices, and React TypeScript web platforms. Led architecture for RAG systems serving 2M+ monthly queries.

TECHNICAL SKILLS

Languages: Python, TypeScript, JavaScript, SQL, C++

Frontend: React, Next.js, Tailwind CSS, Redux Toolkit, HTML5/CSS3

Backend & APIs: FastAPI, Django, Node.js, RESTful APIs, gRPC, Celery

AI & LLMs: Large Language Models (LLM), RAG, LangChain, LlamaIndex, OpenAI, Hugging Face, PyTorch, Vector Databases

Databases: PostgreSQL, Redis, MongoDB, SQLite

Cloud & DevOps: AWS (EC2, S3, RDS), Docker, Kubernetes, CI/CD, GitHub Actions

PROFESSIONAL EXPERIENCE

Senior Full-Stack AI Engineer | NeuralScale AI (2021 - Present)

- Architected and shipped an enterprise GenAI platform using FastAPI, Python, and React TypeScript, reducing latency by 45%.
- Implemented hybrid vector search RAG pipeline utilizing PostgreSQL and OpenAI embeddings.
- Containerized microservices using Docker and orchestrated deployments on AWS ECS and Kubernetes.

Full-Stack Developer | Apex Software Labs (2018 - 2021)

- Developed responsive web applications in React and FastAPI backed by PostgreSQL.
- Designed REST APIs handling 5,000+ requests per second with Redis caching.

EDUCATION

Master of Science (M.S.) in Computer Science - Stanford University (2018)

Bachelor of Science (B.S.) in Computer Science - UC Berkeley (2016)

NOTABLE PROJECTS

PromptFlow-Agent: Open-source framework for structured LLM evaluation and schema validation (Python, FastAPI, React, PyTorch).

Multi-Tenant RAG Engine: Scalable enterprise search engine with hybrid BM25 and vector embeddings (Python, PostgreSQL, Docker, AWS).

CERTIFICATIONS

AWS Certified Solutions Architect - Associate